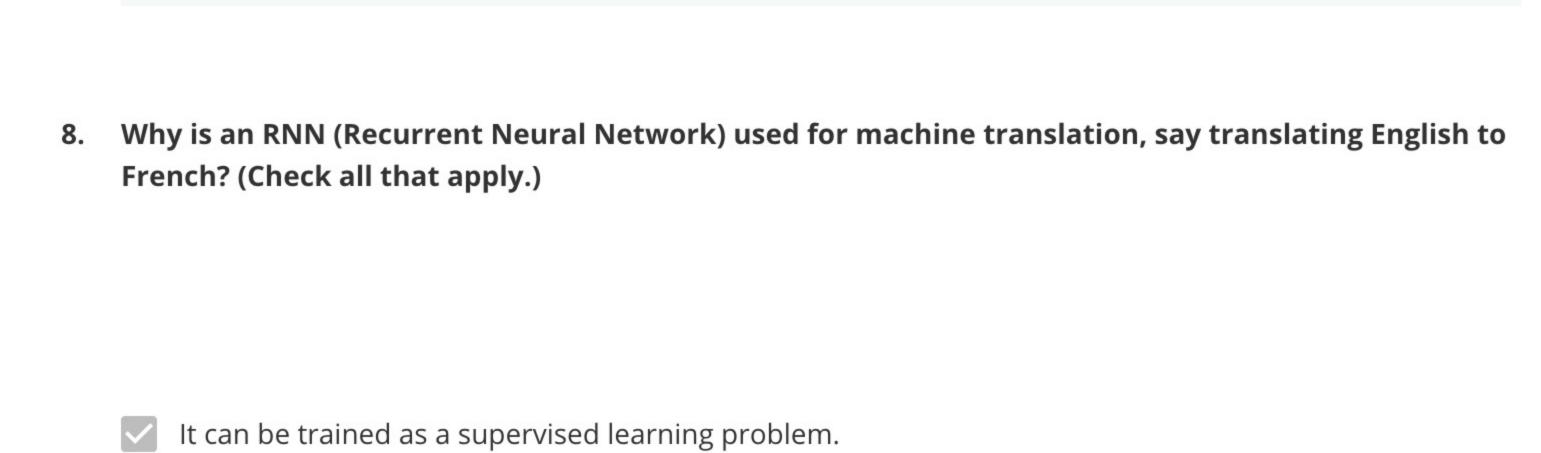
Idea Experiment Being able to try out ideas quickly allows deep learning engineers to iterate more quickly. ✓ Correct Yes, as discussed in Lecture 4. Faster computation can help speed up how long a team takes to iterate to a good idea. ✓ Correct Yes, as discussed in Lecture 4. without changing the CPU/GPU hardware). Correct Yes. For example, we discussed how switching from sigmoid to ReLU activation functions allows faster training. 4. When an experienced deep learning engineer works on a new problem, they can usually use insight 1 / 1 point from previous problems to train a good model on the first try, without needing to iterate multiple times through different models. True/False? True False Correct Yes. Finding the characteristics of a model is key to have good performance. Although Which one of these plots represents a ReLU activation function? 1/1 point Figure 1: Figure 3: Figure 4: ✓ Correct Correct! This is the ReLU activation function, the most used in neural networks. Images for cat recognition is an example of "structured" data, because it is represented as a 1/1 point structured array in a computer. True/False? True False ✓ Correct



A demographic dataset with statistics on different cities' population, GDP per capita, economic

1/1 point

1/1 point

1 / 1 point

Large Neural

Small Neural Network

growth is an example of "structured" data by opposition to image, audio or text datasets.

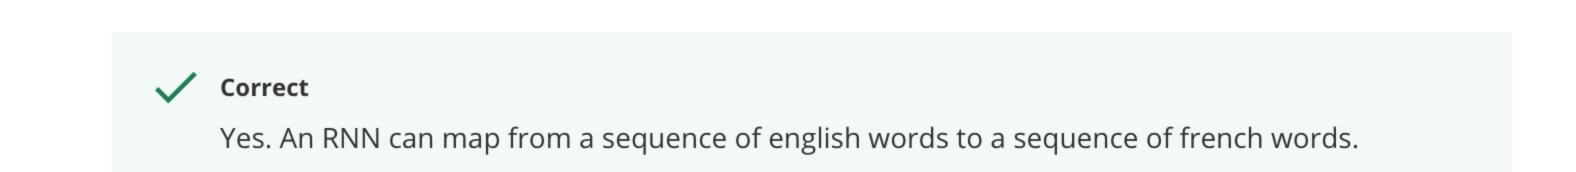
Yes. Images for cat recognition is an example of "unstructured" data.

True

False

✓ Correct

(y-axis) represent?



In this diagram which we hand-drew in lecture, what do the horizontal axis (x-axis) and vertical axis

It is strictly more powerful than a Convolutional Neural Network (CNN).

It is applicable when the input/output is a sequence (e.g., a sequence of words).

RNNs represent the recurrent process of Idea->Code->Experiment->Idea->....





• y-axis is the size of the model you train. x-axis is the input to the algorithm • y-axis is outputs.

• x-axis is the performance of the algorithm

• y-axis (vertical axis) is the amount of data.

x-axis is the amount of data

• y-axis (vertical axis) is the performance of the algorithm.

- 10. Assuming the trends described in the previous question's figure are accurate (and hoping you got the axis labels right), which of the following are true? (Check all that apply.)
 - Increasing the size of a neural network generally does not hurt an algorithm's performance, and it may help significantly.

✓ Correct

- Yes. According to the trends in the figure above, big networks usually perform better than small networks.
- Decreasing the size of a neural network generally does not hurt an algorithm's performance, and it may help significantly. Decreasing the training set size generally does not hurt an algorithm's performance, and it may help
- Increasing the training set size generally does not hurt an algorithm's performance, and it may help
- significantly. significantly.