

Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE
100%


Sequence models & Attention mechanism

LATEST SUBMISSION GRADE

100%

1. Consider using this encoder-decoder model for machine translation.

1 / 1 point



This model is a “conditional language model” in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence x .

- ☐ True
- ☒ False

Correct

2. In beam search, if you increase the beam width B , which of the following would you expect to be true? Check all that apply.

1 / 1 point

- ☒ Beam search will run more slowly.

Correct

Correct

- ☒ Beam search will generally find better solutions (i.e. do a better job maximizing $P(y \mid x)$)

Correct

- ☐ Beam search will converge after fewer steps.

3. In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations.

1 / 1 point

- ☒ True
- ☐ False

4. Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip x to a text transcript y . Your algorithm uses beam search to try to find the value of y that maximizes $P(y \mid x)$.

1 / 1 point

On a dev set example, given an input audio clip, your algorithm outputs the transcript $\hat{y} = \text{“I’m building an A Eye system in Silly con Valley.”}$, whereas a human gives a much superior transcript $y^* = \text{“I’m building an AI system in Silicon Valley.”}$

According to your model,

$$P(\hat{y} \mid x) = 1.09 * 10^{-7}$$

$$P(y^* \mid x) = 7.21 * 10^{-8}$$

Would you expect increasing the beam width B to help correct this example?

- ☒ No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.
- ☐ No, because $P(y^* \mid x) < P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm
- ☐ Yes, because $P(y^* \mid x) < P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

Correct

5. Continuing the example from Q4, suppose you work on your algorithm for a few more weeks, and now find that for the vast majority of examples on which your algorithm makes a mistake, $P(y^* \mid x) > P(\hat{y} \mid x)$. This suggest you should focus your attention on improving the search algorithm.


1 / 1 point

- ☒ True.
- ☐ False.

Correct

6. Consider the attention model for machine translation.

1 / 1 point



Further, here is the formula for $\alpha^{<t,t'>}$.

Which of the following statements about $\alpha^{<t,t'>}$ are true? Check all that apply.

- ☒ We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t>}$. (Note the indices in the superscripts.)

Correct

- ☐ We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t>}$ that are highly relevant to the value the network should output for $y^{<t'>}$. (Note the indices in the superscripts.)

- ☐ $\sum_t \alpha^{<t,t'>} = 1$ (Note the summation is over t .)

- ☒ $\sum_{t'} \alpha^{<t,t'>} = 1$ (Note the summation is over t' .)

7. The network learns where to “pay attention” by learning the values $e^{<t,t'>}$, which are computed using a small neural network:

1 / 1 point

We can’t replace $s^{<t-1>}$ with $s^{<t>}$ as an input to this neural network. This is because $s^{<t>}$ depends on $\alpha^{<t,t'>}$ which in turn depends on $e^{<t,t'>}$; so at the time we need to evaluate this network, we haven’t computed $s^{<t>}$ yet.

- ☒ True
- ☐ False

Correct

8. Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism), we expect the attention model to have the greatest advantage when:

1 / 1 point

- ☒ The input sequence length T_x is large.

Correct

9. Under the CTC model, identical repeated characters not separated by the “blank” character ($_$) are “collapsed” into a single character.

1 / 1 point

- ☐ cokbok
- ☒ cookbook
- ☐ cook book
- ☐ coookkboooooookkk

Correct

10. In trigger word detection, $x^{<t>}$ is:

1 / 1 point

- ☒ Features of the audio (such as spectrogram features) at time t .
- ☐ The t -th input word, represented as either a one-hot vector or a word embedding.
- ☐ Whether the trigger word is being said at time t .
- ☐ Whether someone has just finished saying the trigger word at time t .